

DSC630-T301 Predictive Analytics (2243-1)_week4_Samanta_Rajib

January 8, 2024

0.1 Class : DSC630-T301 Predictive Analytics (2243-1)

0.2 Name : Rajib Samanta

0.2.1 Assignment 4.2 : Week 4

You will be using the dataset `als_data.csv` to apply clustering methods for this assignment. This data gives anonymized data on ALS patients. With this data, complete the following steps: 1. Remove any data that is not relevant to the patient's ALS condition. 2. Apply a standard scalar to the data. 3. Create a plot of the cluster silhouette score versus the number of clusters in a K-means cluster. 4. Use the plot created in (3) to choose on optimal number of clusters for K-means. Justify your choice. 5. Fit a K-means model to the data with the optimal number of clusters chosen in part (4). 6. Fit a PCA transformation with two features to the scaled data. 7. Make a scatterplot the PCA transformed data coloring each point by its cluster value. 8. Summarize your results and make a conclusion. You can use R or Python to complete this assignment. Submit your code and output to the submission link. Make sure to add comments to all of your code and to document your steps, process, and analysis.

```
[36]: # Load the Libraries
import os
import pandas as pd
import matplotlib.pyplot as plt
import %matplotlib inline
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.decomposition import PCA
```

```
[37]: # 1. Load the dataset as a Pandas data frame.
# 2. Display the first ten rows of data.
# Read anonymized data on ALS patients data file ('als_data.csv') from local:
directory = '/Users/rajibsamanta/Documents/Rajib/College/Sem 7 Winter 2023/
↳Week4'
# Set the working directory
```

```

os.chdir(directory)
print(os.getcwd())
dataset1_csv = pd.read_csv("als_data.csv")
dataset1_csv.head(5)
# Display the DataFrame 10 rows

```

/Users/rajibsamanta/Documents/Rajib/College/Sem 7 Winter 2023/Week4

```

[37]:
  ID  Age_mean  Albumin_max  Albumin_median  Albumin_min  Albumin_range  \
0   1         65         57.0           40.5          38.0         0.066202
1   2         48         45.0           41.0          39.0         0.010453
2   3         38         50.0           47.0          45.0         0.008929
3   4         63         47.0           44.0          41.0         0.012111
4   5         63         47.0           45.5          42.0         0.008292

  ALSFRS_slope  ALSFRS_Total_max  ALSFRS_Total_median  ALSFRS_Total_min  ...  \
0    -0.965608                30                28.0                22  ...
1    -0.921717                37                33.0                21  ...
2    -0.914787                24                14.0                10  ...
3    -0.598361                30                29.0                24  ...
4    -0.444039                32                27.5                20  ...

  Sodium_min  Sodium_range  SubjectID  trunk_max  trunk_median  trunk_min  \
0        143.0        0.017422         533         8           7.0           7
1        136.0        0.010453         649         8           7.0           5
2        140.0        0.008929        1234         5           0.0           0
3        138.0        0.012469        2492         5           5.0           3
4        138.0        0.008292        2956         6           4.0           1

  trunk_range  Urine.Ph_max  Urine.Ph_median  Urine.Ph_min
0    0.002646           6.0           6.0           6.0
1    0.005386           7.0           5.0           5.0
2    0.008929           6.0           5.0           5.0
3    0.004988           7.0           6.0           5.0
4    0.008489           6.0           5.0           5.0

```

[5 rows x 101 columns]

1. Remove any data that is not relevant to the patient's ALS condition.

```

[38]: # 1. Remove any data that is not relevant to the patient's ALS condition.
# Identify and remove irrelevant columns
columns_to_remove = ['ID', 'SubjectID'] # Replace with the actual column names_
↳ to be removed
dataset1_csv = dataset1_csv.drop(columns=columns_to_remove, errors='ignore')

# Display the DataFrame after removing irrelevant data
print("\nDataFrame after removing irrelevant data:")

```

```
print(dataset1_csv.head())
```

DataFrame after removing irrelevant data:

	Age_mean	Albumin_max	Albumin_median	Albumin_min	Albumin_range	\
0	65	57.0	40.5	38.0	0.066202	
1	48	45.0	41.0	39.0	0.010453	
2	38	50.0	47.0	45.0	0.008929	
3	63	47.0	44.0	41.0	0.012111	
4	63	47.0	45.5	42.0	0.008292	

	ALSFRS_slope	ALSFRS_Total_max	ALSFRS_Total_median	ALSFRS_Total_min	\
0	-0.965608	30	28.0	22	
1	-0.921717	37	33.0	21	
2	-0.914787	24	14.0	10	
3	-0.598361	30	29.0	24	
4	-0.444039	32	27.5	20	

	ALSFRS_Total_range	...	Sodium_median	Sodium_min	Sodium_range	\
0	0.021164	...	145.5	143.0	0.017422	
1	0.028725	...	138.0	136.0	0.010453	
2	0.025000	...	143.0	140.0	0.008929	
3	0.014963	...	139.0	138.0	0.012469	
4	0.020374	...	140.0	138.0	0.008292	

	trunk_max	trunk_median	trunk_min	trunk_range	Urine.Ph_max	\
0	8	7.0	7	0.002646	6.0	
1	8	7.0	5	0.005386	7.0	
2	5	0.0	0	0.008929	6.0	
3	5	5.0	3	0.004988	7.0	
4	6	4.0	1	0.008489	6.0	

	Urine.Ph_median	Urine.Ph_min
0	6.0	6.0
1	5.0	5.0
2	5.0	5.0
3	6.0	5.0
4	5.0	5.0

[5 rows x 99 columns]

```
[39]: dataset1_csv.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2223 entries, 0 to 2222
Data columns (total 99 columns):
#   Column                                Non-Null Count  Dtype
---  -

```

0	Age_mean	2223 non-null	int64
1	Albumin_max	2223 non-null	float64
2	Albumin_median	2223 non-null	float64
3	Albumin_min	2223 non-null	float64
4	Albumin_range	2223 non-null	float64
5	ALSFRS_slope	2223 non-null	float64
6	ALSFRS_Total_max	2223 non-null	int64
7	ALSFRS_Total_median	2223 non-null	float64
8	ALSFRS_Total_min	2223 non-null	int64
9	ALSFRS_Total_range	2223 non-null	float64
10	ALT.SGPT._max	2223 non-null	float64
11	ALT.SGPT._median	2223 non-null	float64
12	ALT.SGPT._min	2223 non-null	float64
13	ALT.SGPT._range	2223 non-null	float64
14	AST.SGOT._max	2223 non-null	int64
15	AST.SGOT._median	2223 non-null	float64
16	AST.SGOT._min	2223 non-null	float64
17	AST.SGOT._range	2223 non-null	float64
18	Bicarbonate_max	2223 non-null	float64
19	Bicarbonate_median	2223 non-null	float64
20	Bicarbonate_min	2223 non-null	float64
21	Bicarbonate_range	2223 non-null	float64
22	Blood.Urea.Nitrogen..BUN._max	2223 non-null	float64
23	Blood.Urea.Nitrogen..BUN._median	2223 non-null	float64
24	Blood.Urea.Nitrogen..BUN._min	2223 non-null	float64
25	Blood.Urea.Nitrogen..BUN._range	2223 non-null	float64
26	bp_diastolic_max	2223 non-null	int64
27	bp_diastolic_median	2223 non-null	float64
28	bp_diastolic_min	2223 non-null	int64
29	bp_diastolic_range	2223 non-null	float64
30	bp_systolic_max	2223 non-null	int64
31	bp_systolic_median	2223 non-null	float64
32	bp_systolic_min	2223 non-null	int64
33	bp_systolic_range	2223 non-null	float64
34	Calcium_max	2223 non-null	float64
35	Calcium_median	2223 non-null	float64
36	Calcium_min	2223 non-null	float64
37	Calcium_range	2223 non-null	float64
38	Chloride_max	2223 non-null	float64
39	Chloride_median	2223 non-null	float64
40	Chloride_min	2223 non-null	float64
41	Chloride_range	2223 non-null	float64
42	Creatinine_max	2223 non-null	float64
43	Creatinine_median	2223 non-null	float64
44	Creatinine_min	2223 non-null	float64
45	Creatinine_range	2223 non-null	float64
46	Gender_mean	2223 non-null	int64
47	Glucose_max	2223 non-null	float64

48	Glucose_median	2223	non-null	float64
49	Glucose_min	2223	non-null	float64
50	Glucose_range	2223	non-null	float64
51	hands_max	2223	non-null	int64
52	hands_median	2223	non-null	float64
53	hands_min	2223	non-null	int64
54	hands_range	2223	non-null	float64
55	Hematocrit_max	2223	non-null	float64
56	Hematocrit_median	2223	non-null	float64
57	Hematocrit_min	2223	non-null	float64
58	Hematocrit_range	2223	non-null	float64
59	Hemoglobin_max	2223	non-null	float64
60	Hemoglobin_median	2223	non-null	float64
61	Hemoglobin_min	2223	non-null	float64
62	Hemoglobin_range	2223	non-null	float64
63	leg_max	2223	non-null	int64
64	leg_median	2223	non-null	float64
65	leg_min	2223	non-null	int64
66	leg_range	2223	non-null	float64
67	mouth_max	2223	non-null	int64
68	mouth_median	2223	non-null	float64
69	mouth_min	2223	non-null	int64
70	mouth_range	2223	non-null	float64
71	onset_delta_mean	2223	non-null	int64
72	onset_site_mean	2223	non-null	int64
73	Platelets_max	2223	non-null	int64
74	Platelets_median	2223	non-null	float64
75	Platelets_min	2223	non-null	float64
76	Potassium_max	2223	non-null	float64
77	Potassium_median	2223	non-null	float64
78	Potassium_min	2223	non-null	float64
79	Potassium_range	2223	non-null	float64
80	pulse_max	2223	non-null	int64
81	pulse_median	2223	non-null	float64
82	pulse_min	2223	non-null	int64
83	pulse_range	2223	non-null	float64
84	respiratory_max	2223	non-null	int64
85	respiratory_median	2223	non-null	float64
86	respiratory_min	2223	non-null	int64
87	respiratory_range	2223	non-null	float64
88	Sodium_max	2223	non-null	float64
89	Sodium_median	2223	non-null	float64
90	Sodium_min	2223	non-null	float64
91	Sodium_range	2223	non-null	float64
92	trunk_max	2223	non-null	int64
93	trunk_median	2223	non-null	float64
94	trunk_min	2223	non-null	int64
95	trunk_range	2223	non-null	float64

```

96 Urine.Ph_max                2223 non-null    float64
97 Urine.Ph_median             2223 non-null    float64
98 Urine.Ph_min                2223 non-null    float64
dtypes: float64(75), int64(24)
memory usage: 1.7 MB

```

2. Apply a standard scalar to the data.

```

[40]: # 2. Apply a standard scalar to the data.
      # Identify and extract all numeric columns
      numeric_columns = dataset1_csv.select_dtypes(include=['number']).columns

      # Extract data to be scaled
      data_to_scale = dataset1_csv[numeric_columns]

      # Apply Standard Scaling
      scaler = StandardScaler()
      scaled_data = scaler.fit_transform(data_to_scale)

      # Replace the original numeric columns with the scaled data
      dataset1_csv[numeric_columns] = scaled_data

      # Display the DataFrame after applying Standard Scaling
      print("\nDataFrame after applying Standard Scaling to all numeric columns:")
      print(dataset1_csv.head())

```

DataFrame after applying Standard Scaling to all numeric columns:

	Age_mean	Albumin_max	Albumin_median	Albumin_min	Albumin_range \
0	0.917137	3.089417	-1.300781	-0.866550	5.480929
1	-0.574879	-0.622016	-1.112401	-0.553303	-0.347725
2	-1.452535	0.924415	1.148162	1.326179	-0.507103
3	0.741606	-0.003443	0.017880	0.073191	-0.174361
4	0.741606	-0.003443	0.583021	0.386438	-0.573670

	ALSFRS_slope	ALSFRS_Total_max	ALSFRS_Total_median	ALSFRS_Total_min \
0	-0.381450	-0.318520	0.134960	0.247368
1	-0.310907	0.998995	0.888863	0.130839
2	-0.299769	-1.447819	-1.975969	-1.150976
3	0.208801	-0.318520	0.285741	0.480425
4	0.456831	0.057913	0.059570	0.014311

	ALSFRS_Total_range	...	Sodium_median	Sodium_min	Sodium_range \
0	-0.301588	...	2.992342	2.300470	0.260968
1	0.166537	...	-1.198812	-0.278144	-0.489913
2	-0.064100	...	1.595291	1.195350	-0.654169
3	-0.685524	...	-0.639992	0.458603	-0.272701
4	-0.350529	...	-0.081171	0.458603	-0.722774

	trunk_max	trunk_median	trunk_min	trunk_range	Urine.Ph_max \
0	1.028018	0.981832	1.715365	-0.997420	-0.880376
1	1.028018	0.981832	0.867032	-0.388669	0.192665
2	-0.688950	-2.280669	-1.253800	0.398249	-0.880376
3	-0.688950	0.049689	0.018699	-0.477181	0.192665
4	-0.116627	-0.416383	-0.829634	0.300598	-0.880376

	Urine.Ph_median	Urine.Ph_min
0	0.463054	1.868532
1	-1.137208	-0.419151
2	-1.137208	-0.419151
3	0.463054	-0.419151
4	-1.137208	-0.419151

[5 rows x 99 columns]

3. Create a plot of the cluster silhouette score versus the number of clusters in a K-means cluster.

[41]: # 3. Create a plot of the cluster silhouette score versus the number of clusters in a K-means cluster.


```
# Range of cluster numbers to try
cluster_range = range(2, 11) # considering the range 2 to 11

# List to store silhouette scores
silhouette_scores = []

# Perform K-means clustering for different numbers of clusters
for num_clusters in cluster_range:
    kmeans = KMeans(n_clusters=num_clusters, random_state=42)
    cluster_labels = kmeans.fit_predict(scaled_data)

    # Compute silhouette score and append to the list
    silhouette_avg = silhouette_score(scaled_data, cluster_labels)
    silhouette_scores.append(silhouette_avg)

# Plotting the results
plt.plot(cluster_range, silhouette_scores, marker='o')
plt.title('Silhouette Score vs Number of Clusters')
plt.xlabel('Number of Clusters')
plt.ylabel('Silhouette Score')
plt.show()
```



4. Use the plot created in (3) to choose on optimal number of clusters for K-means. **Justify your choice.** -> Using the elbow method for the “elbow” point, which is where the improvement in silhouette score starts to diminish. 7 is the optimal number of cluster.

5. Fit a K-means model to the data with the optimal number of clusters chosen in part (4).

```
[42]: # 5. Fit a K-means model to the data with the optimal number of clusters chosen
      ↪ in part (4).

      # Replace 'optimal_num_clusters' with the number we have identified from the
      ↪ silhouette score plot
      optimal_num_clusters = 7 # Replace with determined optimal number of clusters

      # Fit K-means model with the optimal number of clusters
      kmeans = KMeans(n_clusters=optimal_num_clusters, random_state=42)
      dataset1_csv['cluster_labels'] = kmeans.fit_predict(scaled_data)

      # Display the DataFrame with cluster labels
      print("DataFrame with Cluster Labels:")
```



```
dataset1_csv.head()
```

DataFrame with Cluster Labels:

```
[42]:   Age_mean  Albumin_max  Albumin_median  Albumin_min  Albumin_range  \
0   0.917137    3.089417    -1.300781    -0.866550     5.480929
1  -0.574879    -0.622016    -1.112401    -0.553303    -0.347725
2  -1.452535     0.924415     1.148162     1.326179    -0.507103
3   0.741606    -0.003443     0.017880     0.073191    -0.174361
4   0.741606    -0.003443     0.583021     0.386438    -0.573670

      ALSFRS_slope  ALSFRS_Total_max  ALSFRS_Total_median  ALSFRS_Total_min  \
0    -0.381450      -0.318520           0.134960         0.247368
1    -0.310907       0.998995           0.888863         0.130839
2    -0.299769     -1.447819          -1.975969        -1.150976
3     0.208801     -0.318520           0.285741         0.480425
4     0.456831      0.057913           0.059570         0.014311

      ALSFRS_Total_range  ...  Sodium_min  Sodium_range  trunk_max  trunk_median  \
0    -0.301588  ...    2.300470     0.260968    1.028018     0.981832
1     0.166537  ...   -0.278144    -0.489913    1.028018     0.981832
2    -0.064100  ...    1.195350    -0.654169   -0.688950    -2.280669
3    -0.685524  ...    0.458603    -0.272701   -0.688950     0.049689
4    -0.350529  ...    0.458603    -0.722774   -0.116627    -0.416383

      trunk_min  trunk_range  Urine.Ph_max  Urine.Ph_median  Urine.Ph_min  \
0    1.715365   -0.997420    -0.880376     0.463054     1.868532
1    0.867032   -0.388669     0.192665    -1.137208    -0.419151
2   -1.253800    0.398249    -0.880376    -1.137208    -0.419151
3    0.018699   -0.477181     0.192665     0.463054    -0.419151
4   -0.829634    0.300598    -0.880376    -1.137208    -0.419151

      cluster_labels
0                0
1                6
2                5
3                2
4                6
```

[5 rows x 100 columns]

6. Fit a PCA transformation with two features to the scaled data.

```
[43]: # 6. Fit a PCA transformation with two features to the scaled data.
# Fit PCA transformation with two features
pca = PCA(n_components=2)
pca_result = pca.fit_transform(scaled_data)
```

```
# Create a new DataFrame with the PCA results
pca_df = pd.DataFrame(data=pca_result, columns=['PC1', 'PC2'])

# Display the DataFrame with PCA results
print("DataFrame with PCA Results:")
pca_df.head()
```

DataFrame with PCA Results:


```
[43]:      PC1      PC2
0 -1.426768 -2.320223
1 -1.440195 -4.870052
2  1.617824 -0.430010
3 -1.919997  2.096063
4  0.297701  0.166783
```

7. Make a scatterplot the PCA transformed data coloring each point by its cluster value.

```
[44]: # 7. Make a scatterplot the PCA transformed data coloring each point by its
      ↪ cluster value.

pca_df['Cluster'] = dataset1_csv['cluster_labels']
# Scatterplot with cluster colors
plt.figure(figsize=(10, 6))
for cluster_label in range(optimal_num_clusters):
    cluster_data = pca_df[pca_df['Cluster'] == cluster_label]
    plt.scatter(cluster_data['PC1'], cluster_data['PC2'], label=f'Cluster_{
    ↪{cluster_label}')
```

```
plt.title('Scatterplot of PCA Transformed Data with Cluster Colors')
plt.xlabel('Principal Component 1 (PC1)')
plt.ylabel('Principal Component 2 (PC2)')
plt.legend()
plt.show()
```



```
[45]: # Count of records in each cluster
cluster_counts = dataset1_csv['cluster_labels'].value_counts()

# Display the count of records in each cluster
print("Record Count in Each Cluster:")
print(cluster_counts)
```

Record Count in Each Cluster:

cluster_labels

4 472

2 466

6 419

0 338

5 331

3 171

1 26

Name: count, dtype: int64

8. Summarize your results and make a conclusion. ***The analysis aimed to uncover patterns and relationships within ALS patient data by combining K-means clustering and PCA. The silhouette score plot was used to determine the optimal number of clusters, and a K-means model was applied to assign cluster labels to the data. PCA was then employed to reduce the dimensionality of the data for visualization. The scatterplot of the PCA-transformed data with cluster colors provides insights into the grouping of ALS patients based on the identified features.*

The visual representation aids in understanding potential clusters or patterns that may exist in the data.

Cluster 2 ,4,5,6 -> Majority of the records are under these clusters, and they are well distributed. These clusters data provide the grouping of ALS patients based the diagnosis data. Which will help to do patterns analysis.

Cluster 0 & 3 -> records are spread out across another zone. Looks like more data cleanup required or we are considering unwanted rows/columns which causing the overlapping clustering.

Cluster 1 -> There are only 26 records out of 2.2k records, These may the rare scenarios or outlier.

**Further domain-specific knowledge and additional statistical assessments may be needed to draw meaningful conclusions and insights from the clustering and dimensionality reduction results

[]:

